



Update Consistency for Wait-free Concurrent Objects

Matthieu Perrin Achour Mostefaoui Claude Jard

LINA, University of Nantes

May 26th, 2015

Architecture of a distributed application

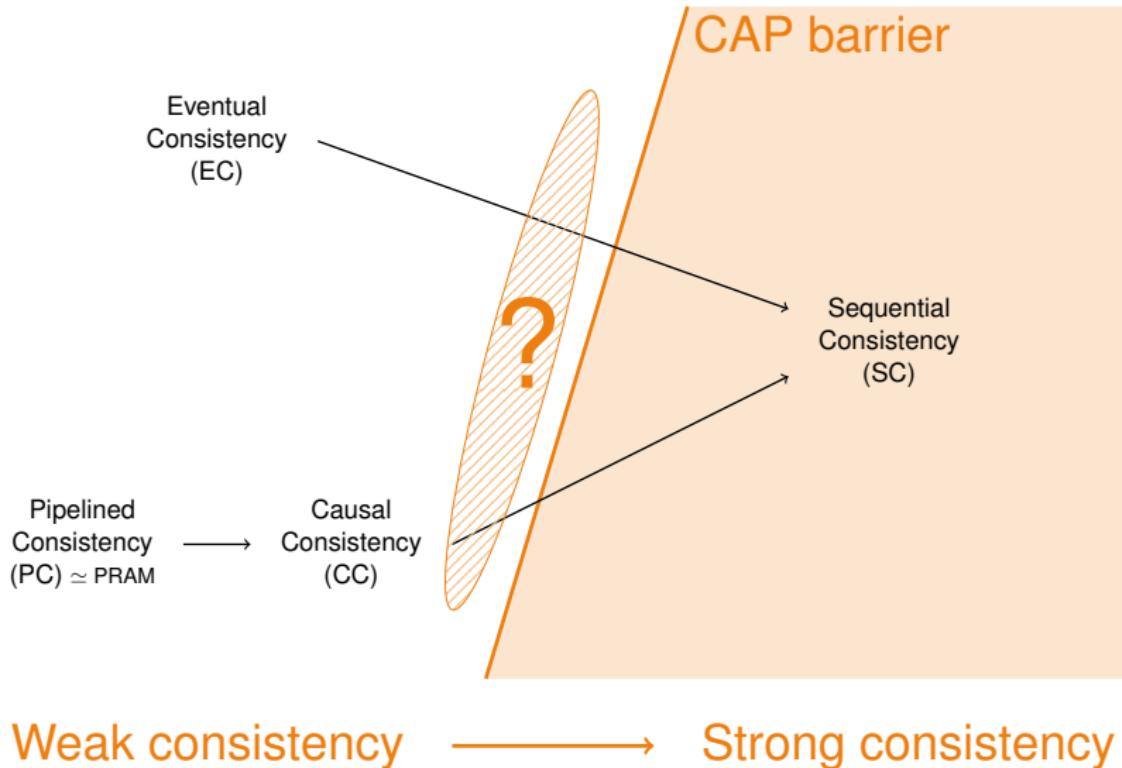


Shared Objects in Distributed Systems

- ▶ Strong Consistency
- ▶ Availability
- ▶ Partition tolerance

Impossible [Gilbert, Lynch, 2002]

A map of weak consistency



Plan

1. Specifying shared objects
2. Towards Pipelined Convergence?
3. Update Consistency

Specifying shared objects

Sequential specification

- ▶ How the object behaves in a sequential environment
 - ▶ Memory
 - ▶ Set
 - ▶ Data base

Consistency criterion

- ▶ How distribution impacts the sequential behaviour
 - ▶ Causal consistency
 - ▶ Eventual consistency
 - ▶ Sequential consistency

Update-Query Abstract Data Type (UQ-ADT)

Transition system

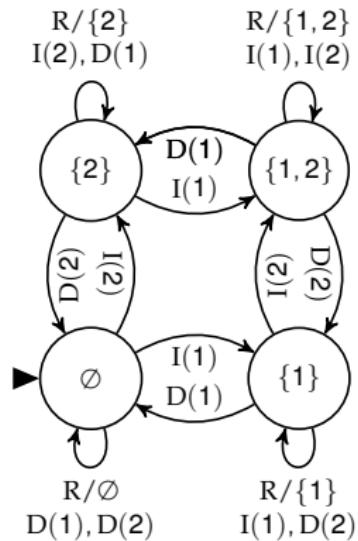
- ▶ Update operations
 - ▶ change the state
 - ▶ have no output
- ▶ Query operation
 - ▶ loop over the states
 - ▶ return a value

Update-Query Abstract Data Type (UQ-ADT)

Transition system

- ▶ Update operations
 - ▶ change the state
 - ▶ have no output
- $\{I(1), I(2), D(1), D(2)\}$
- ▶ Query operation
 - ▶ loop over the states
 - ▶ return a value
- $\{R/\emptyset, R/\{1\}, R/\{2\}, R/\{1, 2\}\}$

E.g. set $\mathcal{S}_{\{1,2\}}$



Update-Query Abstract Data Type (UQ-ADT)

Transition system

- ▶ Update operations
 - ▶ change the state
 - ▶ have no output

$\{I(1), I(2), D(1), D(2)\}$

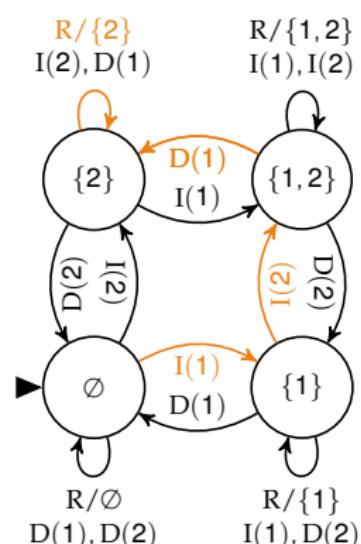
- ▶ Query operation
 - ▶ loop over the states
 - ▶ return a value

$\{R/\emptyset, R/\{1\}, R/\{2\}, R/\{1, 2\}\}$

Sequential specification

- ▶ Set of all the paths

$$I(1)I(2)D(1)R/\{2\}^\omega \in L(\mathcal{S}_{\{1,2\}})$$



Distributed history

Distributed history $H \in \mathbb{H}$

- ▶ Partial order of events, labelled by updates and queries
- ▶ Each event has a finite past

Distributed history

Distributed history $H \in \mathbb{H}$

- ▶ Partial order of events, labelled by updates and queries
- ▶ Each event has a finite past

Wait-free system

- ▶ Asynchronous processes
- ▶ Asynchronous lossless channel
- ▶ Any number of crashes

Admitted history

Events: method calls

Labels: method name and returned value

Order: process order

Consistency criterion

Sequential specification

- $L : \text{UQ-ADT} \rightarrow \mathcal{P}((U \cup Q)^\omega \cup (U \cup Q)^*)$

Consistency criterion

- $C : \text{UQ-ADT} \rightarrow \mathcal{P}(\mathbb{H})$
- E.g. Sequential Consistency
 - $SC : T \mapsto \{H \in \mathbb{H} : \text{lin}(H) \cap L(T) \neq \emptyset\}$
- Strength
 - C_1 stronger than C_2 :
 - $\forall T, C_1(T) \subset C_2(T)$
 - C strong:
 - not all UQ-ADTs can be implemented in a wait-free system

Eventual Consistency

Definition

- If the processes stop updating, they eventually converge

$$EC : T \mapsto \{H \in \mathbb{H} : Write4Ever \vee Converge\}$$

Write4Ever $\equiv |U_H| = \infty$

Converge $\equiv \exists Q' \subset Q_H \text{ with } |Q'| < \infty \text{ and all the queries of } Q_H \setminus Q' \text{ can be done in the same state}$

Eventual Consistency

Definition

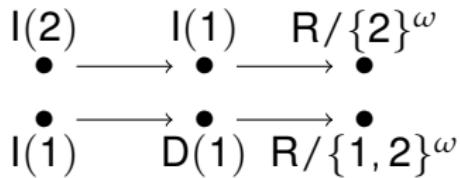
- If the processes stop updating, they eventually converge

$$EC : T \mapsto \{H \in \mathbb{H} : Write4Ever \vee Converge\}$$

Write4Ever $\equiv |U_H| = \infty$

Converge $\equiv \exists Q' \subset Q_H \text{ with } |Q'| < \infty \text{ and all the queries of } Q_H \setminus Q' \text{ can be done in the same state}$

Examples



Eventual Consistency

Definition

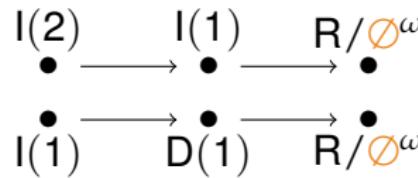
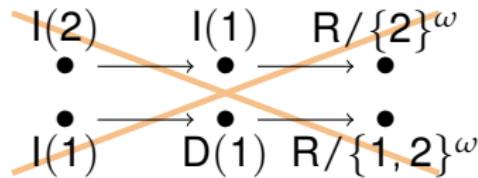
- If the processes stop updating, they eventually converge

$$EC : T \mapsto \{H \in \mathbb{H} : Write4Ever \vee Converge\}$$

Write4Ever $\equiv |U_H| = \infty$

Converge $\equiv \exists Q' \subset Q_H \text{ with } |Q'| < \infty \text{ and all the queries of } Q_H \setminus Q' \text{ can be done in the same state}$

Examples



Pipelined Consistency

Definition

- ▶ For each process, the history could be sequentially consistent if the other processes read correctly

$$PC : T \mapsto \{H \in \mathbb{H} : \forall p, \text{lin}(H_{U_H \cup p}) \cap L(T) \neq \emptyset\}$$

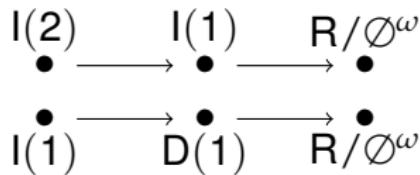
Pipelined Consistency

Definition

- ▶ For each process, the history could be sequentially consistent if the other processes read correctly

$$PC : T \mapsto \{H \in \mathbb{H} : \forall p, \text{lin}(H_{U_H \cup p}) \cap L(T) \neq \emptyset\}$$

Examples



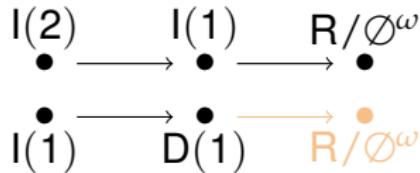
Pipelined Consistency

Definition

- ▶ For each process, the history could be sequentially consistent if the other processes read correctly

$$PC : T \mapsto \{H \in \mathbb{H} : \forall p, \text{lin}(H_{U_H \cup p}) \cap L(T) \neq \emptyset\}$$

Examples



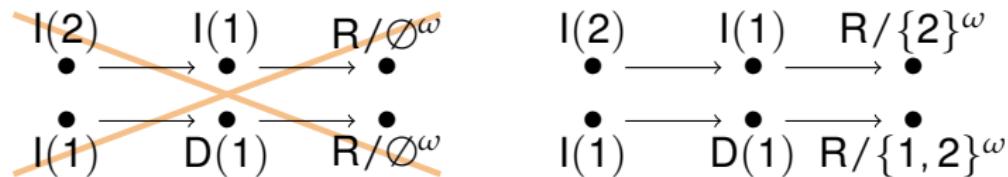
Pipelined Consistency

Definition

- For each process, the history could be sequentially consistent if the other processes read correctly

$$PC : T \mapsto \{H \in \mathbb{H} : \forall p, \text{lin}(H_{U_H \cup p}) \cap L(T) \neq \emptyset\}$$

Examples



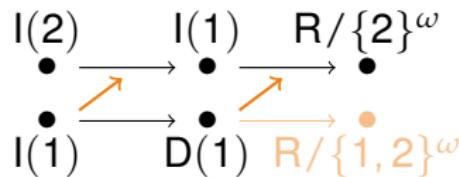
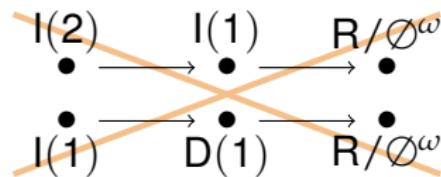
Pipelined Consistency

Definition

- For each process, the history could be sequentially consistent if the other processes read correctly

$$PC : T \mapsto \{H \in \mathbb{H} : \forall p, \text{lin}(H_{U_H \cup p}) \cap L(T) \neq \emptyset\}$$

Examples



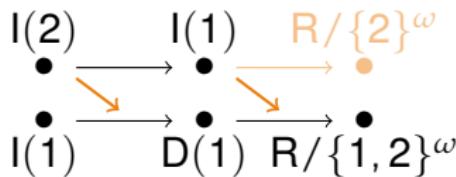
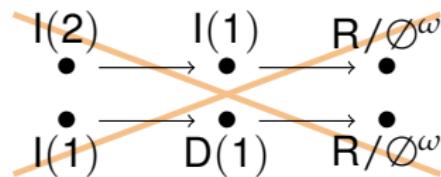
Pipelined Consistency

Definition

- For each process, the history could be sequentially consistent if the other processes read correctly

$$PC : T \mapsto \{H \in \mathbb{H} : \forall p, \text{lin}(H_{U_H \cup p}) \cap L(T) \neq \emptyset\}$$

Examples



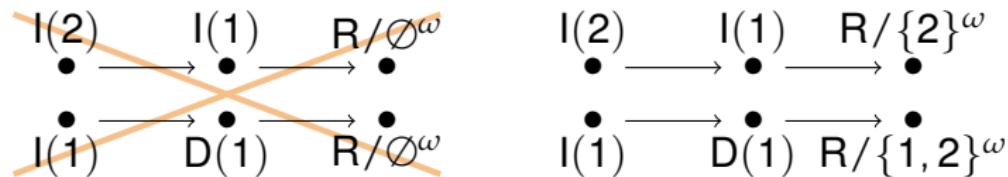
Pipelined Consistency

Definition

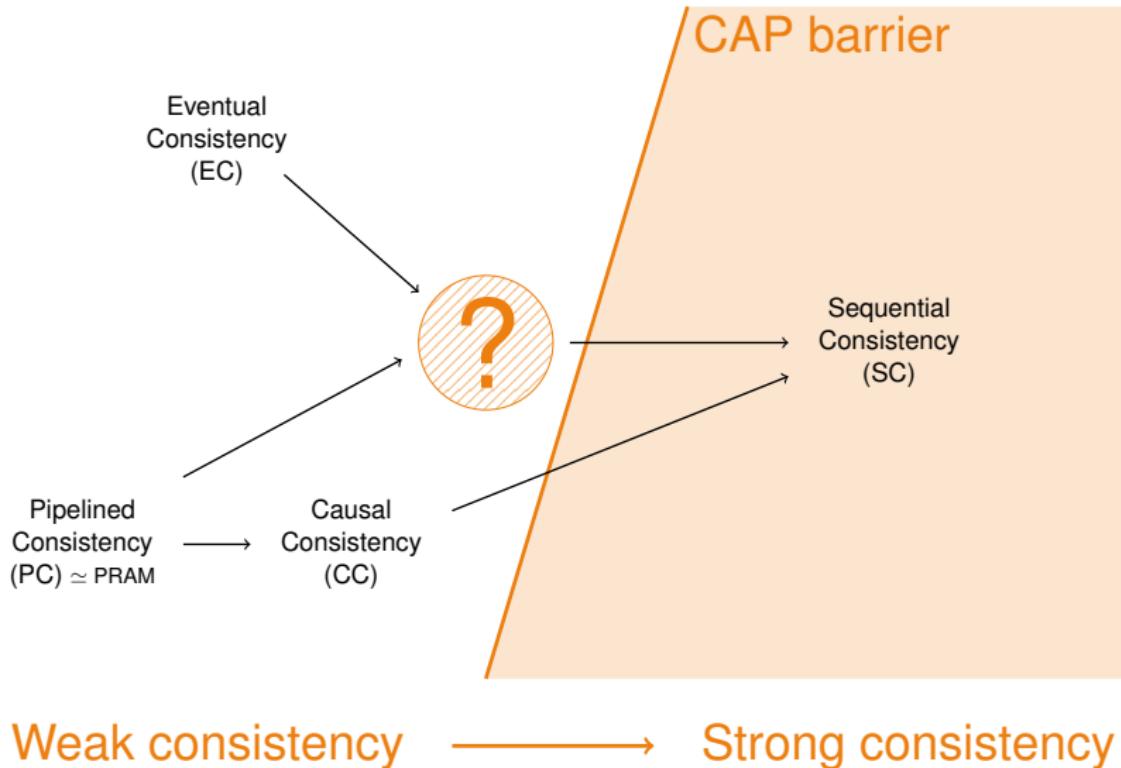
- For each process, the history could be sequentially consistent if the other processes read correctly

$$PC : T \mapsto \{H \in \mathbb{H} : \forall p, \text{lin}(H_{U_H \cup p}) \cap L(T) \neq \emptyset\}$$

Examples



A map of weak consistency



Towards Pipelined Convergence?

Definition

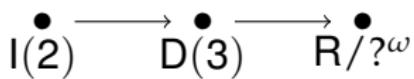
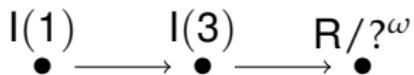
$$EC + PC : T \mapsto EC(T) \cap PC(T)$$

Towards Pipelined Convergence?

Definition

$$EC + PC : T \mapsto EC(T) \cap PC(T)$$

Strength

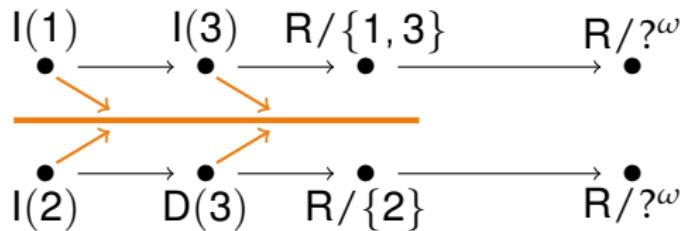


Towards Pipelined Convergence?

Definition

$$EC + PC : T \mapsto EC(T) \cap PC(T)$$

Strength

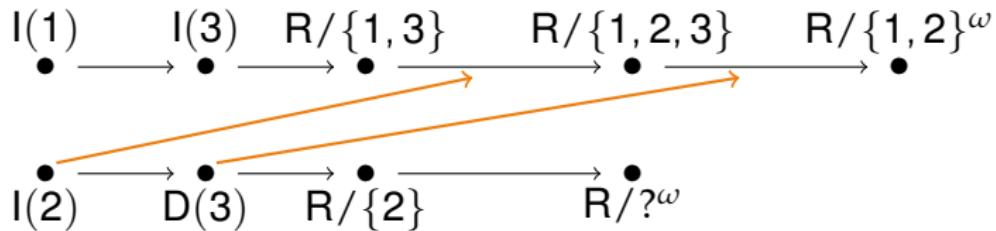


Towards Pipelined Convergence?

Definition

$$EC + PC : T \mapsto EC(T) \cap PC(T)$$

Strength



Towards Pipelined Convergence?

Definition

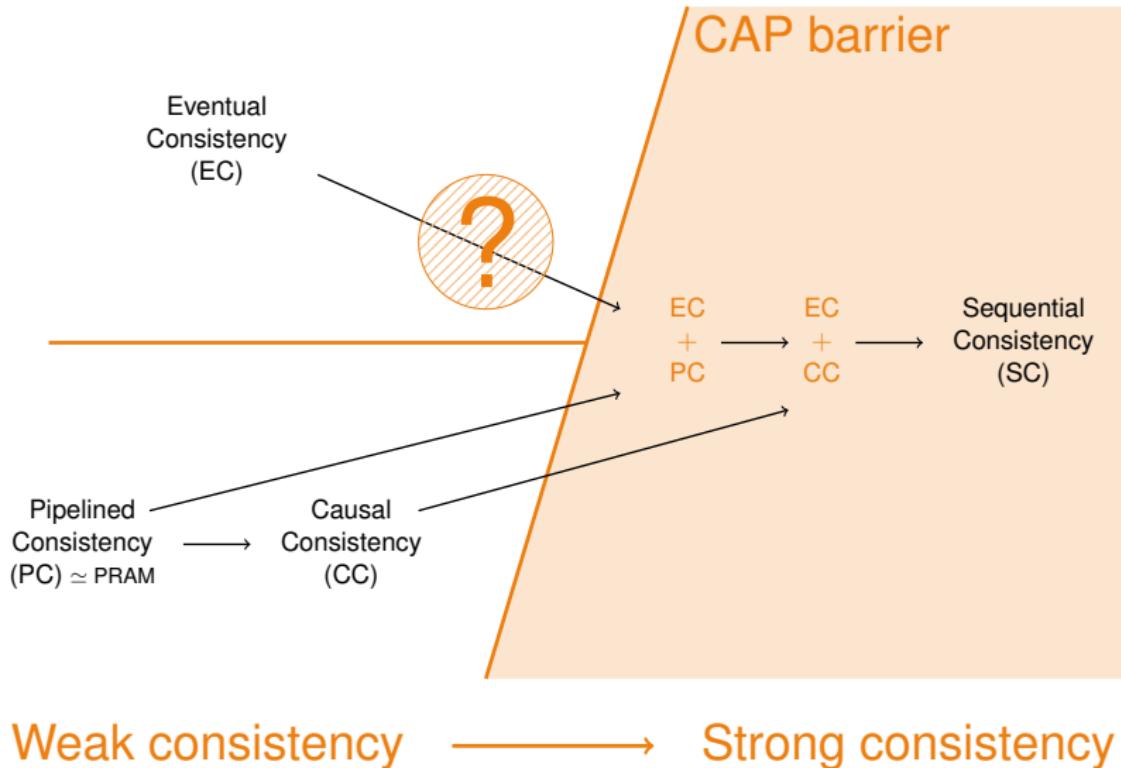
$$EC + PC : T \mapsto EC(T) \cap PC(T)$$

Strength

$$\overset{\bullet}{I(1)} \xrightarrow{\hspace{1cm}} \overset{\bullet}{I(3)} \xrightarrow{\hspace{1cm}} \overset{\bullet}{R/\{1,3\}} \xrightarrow{\hspace{1cm}} \overset{\bullet}{R/\{1,2,3\}} \xrightarrow{\hspace{1cm}} \overset{\bullet}{R/\{1,2\}^\omega}$$

$$\overset{\bullet}{I(2)} \xrightarrow{\hspace{1cm}} \overset{\bullet}{D(3)} \xrightarrow{\hspace{1cm}} \overset{\bullet}{R/\{2\}} \xrightarrow{\hspace{1cm}} \overset{\bullet}{R/\{1,2\}} \xrightarrow{\hspace{1cm}} \overset{\bullet}{R/\{1,2,3\}^\omega}$$

A map of weak consistency



Update Consistency

Definition of Eventual Consistency

$$EC : T \mapsto \{H \in \mathbb{H} : Write4Ever \vee Converge\}$$

$Write4Ever \equiv |U_H| = \infty$

$Converge \equiv \exists Q' \subset Q_H \text{ with } |Q'| < \infty :$

all the queries of $Q_H \setminus Q'$ can be done in the
same state

Update Consistency

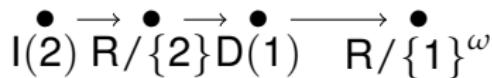
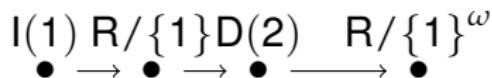
Definition of Update Consistency

$$UC : T \mapsto \{H \in \mathbb{H} : Write4Ever \vee Converge\}$$

Write4Ever $\equiv |U_H| = \infty$

Converge $\equiv \exists Q' \subset Q_H \text{ with } |Q'| < \infty : lin(H_{E \setminus Q'}) \cap L(T) \neq \emptyset$
all the queries of $Q_H \setminus Q'$ can be done in the same state

Examples



Update Consistency

Definition of Update Consistency

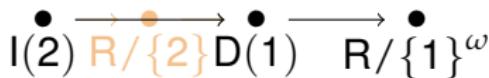
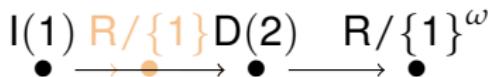
$$UC : T \mapsto \{H \in \mathbb{H} : Write4Ever \vee Converge\}$$

Write4Ever $\equiv |U_H| = \infty$

Converge $\equiv \exists Q' \subset Q_H \text{ with } |Q'| < \infty : lin(H_{E \setminus Q'}) \cap L(T) \neq \emptyset$

all the queries of $Q_H \setminus Q'$ can be done in the same state

Examples



Update Consistency

Definition of Update Consistency

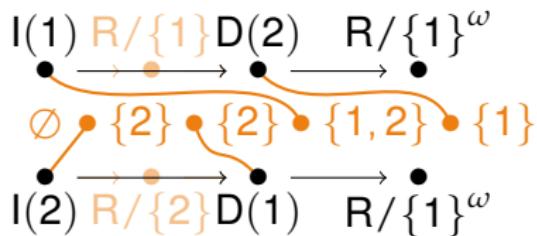
$$UC : T \mapsto \{H \in \mathbb{H} : Write4Ever \vee Converge\}$$

Write4Ever $\equiv |U_H| = \infty$

Converge $\equiv \exists Q' \subset Q_H \text{ with } |Q'| < \infty : lin(H_{E \setminus Q'}) \cap L(T) \neq \emptyset$

all the queries of $Q_H \setminus Q'$ can be done in the same state

Examples



Update Consistency

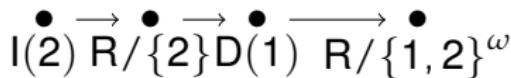
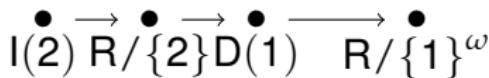
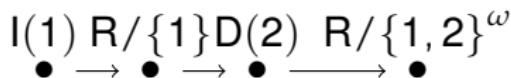
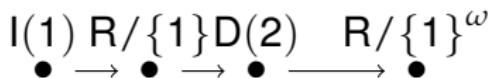
Definition of Update Consistency

$$UC : T \mapsto \{H \in \mathbb{H} : Write4Ever \vee Converge\}$$

Write4Ever $\equiv |U_H| = \infty$

Converge $\equiv \exists Q' \subset Q_H \text{ with } |Q'| < \infty : lin(H_{E \setminus Q'}) \cap L(T) \neq \emptyset$
all the queries of $Q_H \setminus Q'$ can be done in the same state

Examples



Update Consistency

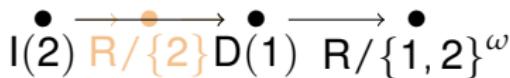
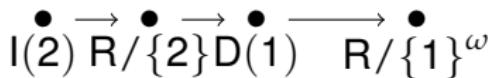
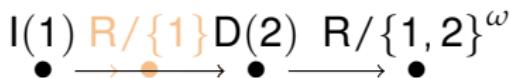
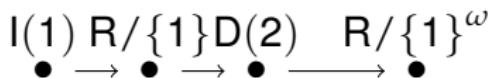
Definition of Update Consistency

$$UC : T \mapsto \{H \in \mathbb{H} : Write4Ever \vee Converge\}$$

Write4Ever $\equiv |U_H| = \infty$

Converge $\equiv \exists Q' \subset Q_H \text{ with } |Q'| < \infty : lin(H_{E \setminus Q'}) \cap L(T) \neq \emptyset$
all the queries of $Q_H \setminus Q'$ can be done in the same state

Examples



Update Consistency

Definition of Update Consistency

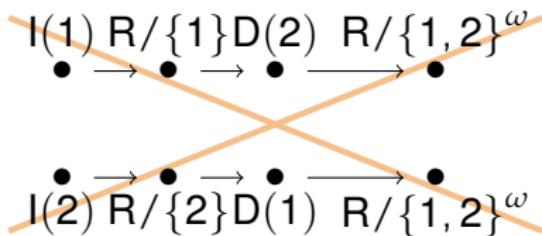
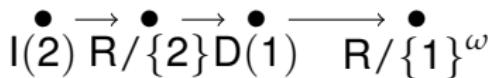
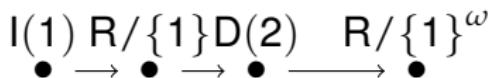
$$UC : T \mapsto \{H \in \mathbb{H} : Write4Ever \vee Converge\}$$

Write4Ever $\equiv |U_H| = \infty$

Converge $\equiv \exists Q' \subset Q_H \text{ with } |Q'| < \infty : lin(H_{E \setminus Q'}) \cap L(T) \neq \emptyset$

all the queries of $Q_H \setminus Q'$ can be done in the same state

Examples

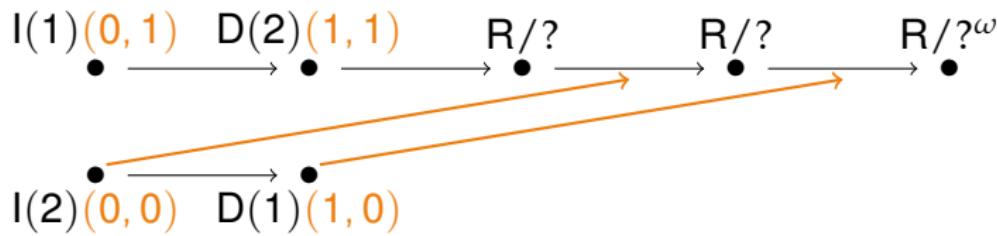


Implementation

Idea

- ▶ Choose a total order *a priori*
- ▶ Update: notify everyone
- ▶ Store all the updates
- ▶ Query: execute the updates in the chosen order

Example

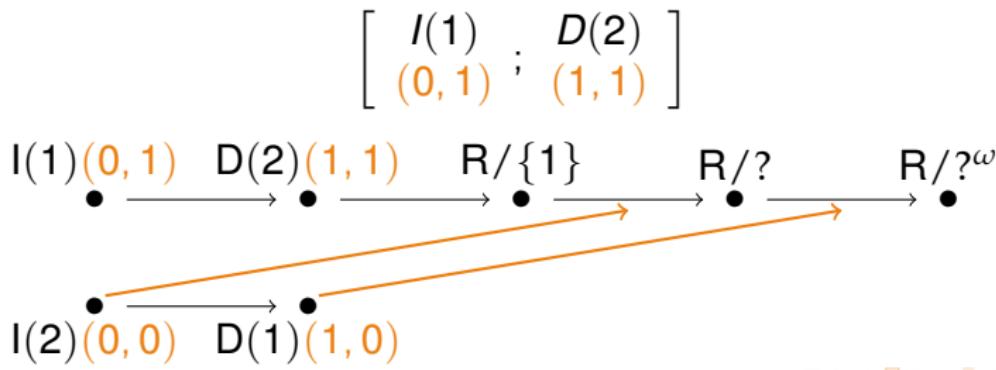


Implementation

Idea

- ▶ Choose a total order *a priori*
- ▶ Update: notify everyone
- ▶ Store all the updates
- ▶ Query: execute the updates in the chosen order

Example

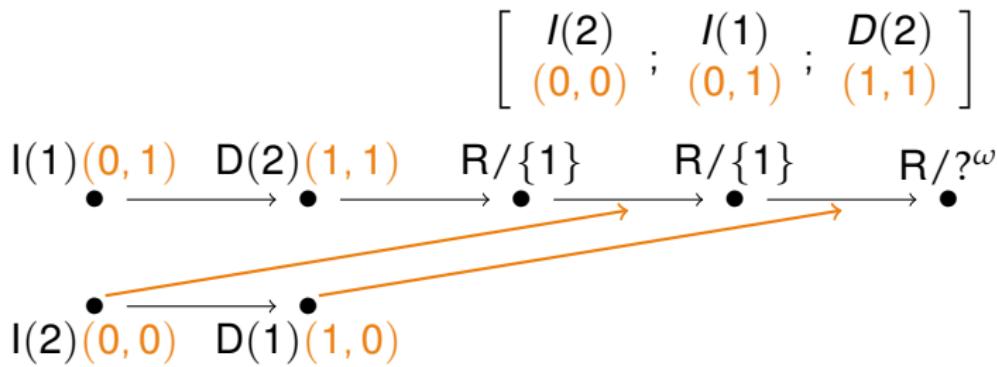


Implementation

Idea

- ▶ Choose a total order *a priori*
- ▶ Update: notify everyone
- ▶ Store all the updates
- ▶ Query: execute the updates in the chosen order

Example

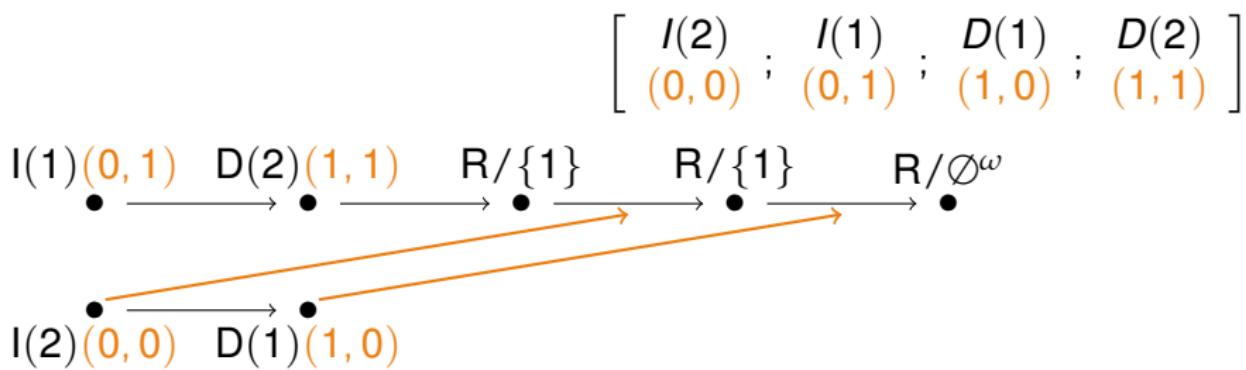


Implementation

Idea

- ▶ Choose a total order *a priori*
- ▶ Update: notify everyone
- ▶ Store all the updates
- ▶ Query: execute the updates in the chosen order

Example

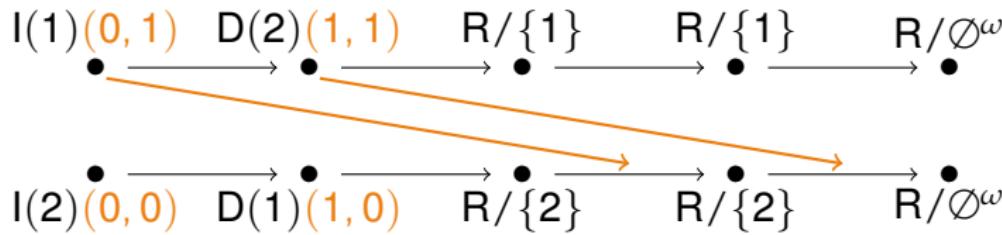


Implementation

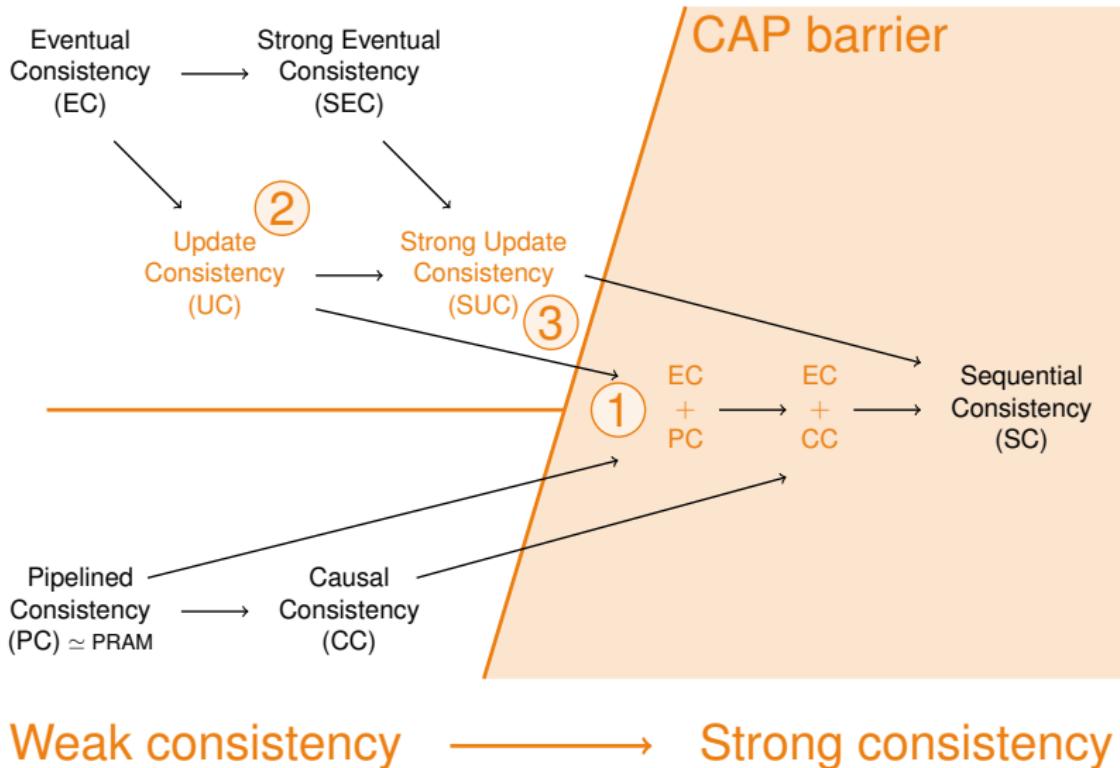
Idea

- ▶ Choose a total order *a priori*
- ▶ Update: notify everyone
- ▶ Store all the updates
- ▶ Query: execute the updates in the chosen order

Example



A map of weak consistency



Future work

Memory-aware implementation

- ▶ So far, only a theoretical proof
- ▶ Unbounded memory/computation

Topology of weak consistency space

- ▶ There is no maximum
- ▶ Are there maximal elements?

Add some causality

- ▶ Causal consistency only defined for memory
- ▶ Exploring the area between SUC and CC+EC